Data Intake Report

Name: Hate Speech detection using Transformers
Report date: 09/01/23
Internship Batch: LISUM25
Version: 1.0
Data intake by: Faris Chaudhry
Data intake reviewer: N/A
Data storage location: https://github.com/farischaudhry/twitter-sentiment-analysis

Tabular data details:

test_tweets.csv

Total number of observations	16130
Total number of files	1
Total number of features	2
Base format of the file	.csv
Size of the data	1.56 MB

train.tweets.csv

Total number of observations	29530
Total number of files	1
Total number of features	3
Base format of the file	.csv
Size of the data	2.96 MB

Proposed Approach:

Data in the training file is labelled whereas in the test file it is not. The ID gives a primary key for the data. Dedup validation can be performed on both files by removing any duplicated tweets or, where two different tweets have the same ID, changing the ID of one.
Duplication, null, and incorrect keys can easily be checked beforehand on Kaggle:

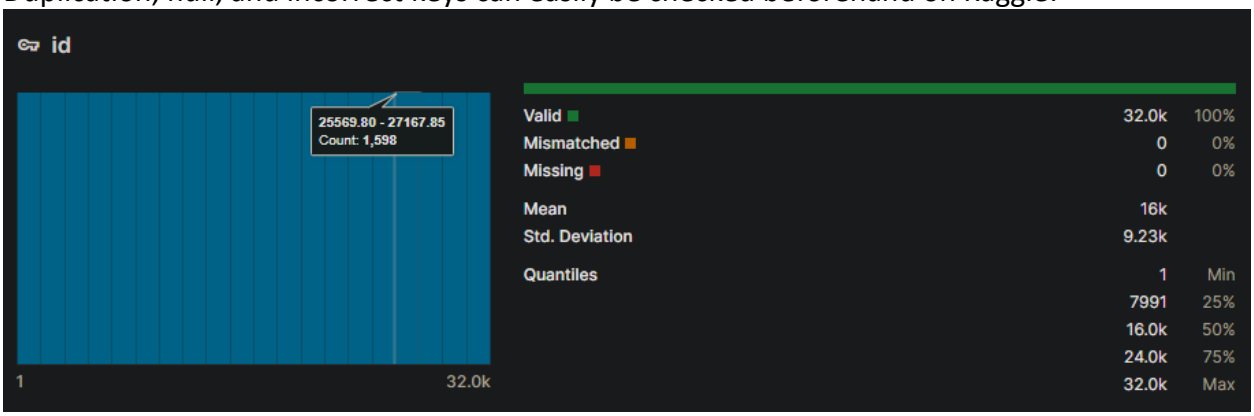

Assumptions:

- The data is derived from real tweets.
- The training data is labelled correctly.
- The training and test data are from the same domain.
- The amount of hate speech compared to non-hate speech reflects the proportion on the platform (see below). Approximately 6% of the labels are classified as hate speech.
- Tweets are below the limit of 200 characters and are formatted in the same way.